

Systems Software HS15

Lab Exercises

Claudio Mura
claudio@ifi.uzh.ch

Practical notes on UNIX Processes

Processes and Identifiers

- a process is a program in execution
 - the “dynamic version” of a program
- the OS keeps some information for each process
 - process id (*pid*), state, parent pid, children, open files, ...
- process id is a fundamental piece of information
 - represented with type `pid_t`
 - `getpid()` function to retrieve it

```
pid_t my_pid = getpid();  
std::cout << "Pid of current process:" << my_pid << "\n";  
  
pid_t parent_pid = getpid();  
std::cout << "Pid of parent process:" << parent_pid << "\n";
```


Creating a Process: `fork()`

- creates a new process (*child*) from an existing process (*parent*)
 - the child is a copy of the parent
 - content of text, stack and data segments is copied
 - some major differences (eg: pid, parent pid)
- returns two times
 - once in the child, once in the parent
- returns a `pid_t` variable containing:
 - the pid of the child, in the parent process
 - the value 0, in the child process
 - the value -1, if the process could not be created

Waiting for Child Process

- after `fork()` returns, the parent and the child process execute concurrently
- when a child completes its execution, it remains in a *zombie* state
 - finished execution, inactive, but its resources can not be deallocated until the parent *waits* for him
- the parent process should call the `wait()` function
 - suspend its execution until a child terminates
 - can wait for a specific child using `waitpid()`
 - returns immediately if a child has already terminated

Checking Child Termination

- both `wait` and `waitpid` take an `int*` `status` argument
 - pointer to a previously declared integer variable
 - information about status of terminated process
- a set of macros are available to check how the process has terminated
 - `WIFEXITED`, `WIFSTOPPED`, `WIFSIGNALED`, ...
 - read the man pages for other useful macros

Example of `wait()`

```
pid_t pid = fork();

if( pid == -1 ) {
    std::cerr << "Error on fork() " << std::endl;
    exit( 1 );
}

if( pid == 0 ) {
    pid_t my_pid = getpid();
    std::cout << "Child process, pid = " << my_pid << std::endl;
    exit( EXIT_SUCCESS );
}
else {
    int status;
    wait( &status );

    if( WIFEXITED( status ) ) {
        std::cout << "Child terminated normally" << std::endl;
        return 0;
    }
    else {
        // handle error
    }
}
```


Changing Process Image

- a process is a program in execution
- using the `exec()` family of functions it is possible to change the program run by a process
 - common practice: use `fork`, then call `exec` in the child
- the text, stack and data segment of the process are replaced
- you must provide to `exec()`:
 - fullpath of executable of the new program
 - input arguments for the new program
 - by convention, the first is the program name
- Note: if successful, `exec` does not return!

Example of exec()

```
#include <iostream>
#include <unistd.h>

int main ( int argc, char* argv[] )
{
    execl ( "/bin/echo",
            "echo",
            "Message to be displayed",
            NULL );

    std::cerr << "Error: exec has returned!" << std::endl;
    return 1;
}
```


Exercise 2

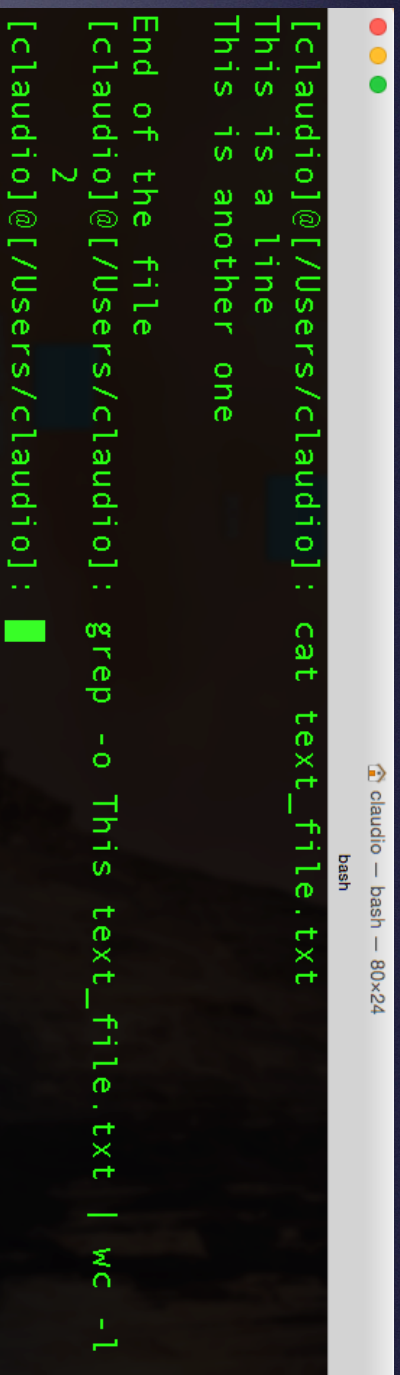
Process creation and management

Parallel string search in text files

- Write a multi-process program that, given a string (*pattern*) and a list of text files, counts the occurrences of the pattern inside the text files
 - pattern and filenames passed as command line arguments
- You should create multiple processes:
 - one child process for every text file
 - each child performs the counting on one file, then outputs the result to a text file
 - use `fork()` from the main process to create the children
 - main process reads output files, computes overall number of occurrences and prints it to standard output

Counting occurrences in a file

- Don't code it yourself! You must use `exec*` to execute an instance of a shell that does it for you
 - using the following command: `grep -o PATTERN FILENAME | wc -l`



```
claudio — bash — 80x24
bash
[claudio]@[/Users/claudio]: cat text_file.txt
This is a line
This is another one
End of the file
[claudio]@[/Users/claudio]: grep -o This text_file.txt | wc -l
2
[claudio]@[/Users/claudio]:
```

- `add > OUTPUT_FILENAME` to write the result to a file
- the output file with the result should be called `result-PID.txt`, with `PID` being the process id of the child process

Counting occurrences in a file

- With `exec` (example):

```
execl ( "/bin/sh" , "/bin/sh" ,  
        "-c" ,  
        "grep -o This text_file.txt | wc -l > out.txt" ,  
        ( char* ) 0 ) ;
```


Assembling the results

- Task of the main process
 - open the output files generated by the child processes
 - all children must have generated their result file - so they must have terminated, and successfully!
 - wait for the processes and check termination status before reading the files
 - from each output file, read the *partial* occurrences count
 - i.e., the number of occurrences found by the child process that generated that output file
 - sum up the partial occurrences and print the result to standard output

Additional comments

- Handle the following abnormal situations appropriately
 - insufficient number of command-line arguments
 - `fork()` or `exec()` fail
 - a child returns abnormally
 - errors while opening file
- A draft of the solution is provided together with the description
 - you may use it at your convenience