

Systems Software HS15

Lab Exercises

Claudio Mura
claudio@ifi.uzh.ch

Practical notes on parallel programming using OpenMP

OpenMP

- High-level API for parallelizing C/C++ and Fortran code
 - based on `#pragma` compiler directives
 - includes a run-time library
- How does OpenMP work?
 - place directives before blocks to be parallelized
 - compiler processes these directives and generates multi-thread code
- Easy to parallelize existing code!

OpenMP Directives

- An OpenMP directive has the following structure:

```
#pragma omp directive [ clause list ]
```

- Example:

```
int a[ SIZE ];  
int i;  
#pragma omp parallel for shared( a ) private( i )  
for( i=0; i<SIZE; i++ )  
{  
    a[ i ] = pow( a[ i ], 2.0f );  
}
```


parallel Directive

- The most important directive is `parallel` :

```
#pragma omp parallel [ clause list ]  
{  
    // structured block  
}
```

- when the main thread encounters this directives, it becomes the *master* and generates a group of threads
- You will mostly need it to parallelize `for` loops
 - syntax shortcut: `#pragma omp parallel for`

parallel Directive: Clauses

- The `parallel` directive supports a number of clauses:

- conditional parallelization:

`#pragma omp parallel for if(scalar expression)`

- degree of concurrency

`#pragma omp parallel for num_threads(integer)`

- data sharing attribute clauses

- `#pragma omp parallel for private(variable list)`

variables are local to each thread (each thread has a copy)

- `#pragma omp parallel for shared(variable list)`

variables are shared among all the threads (only one copy!)

- `#pragma omp parallel for firstprivate(variable list)`

each thread has a copy + all the copies are initialized

- many other

Scheduling Policies for `omp parallel for`

- The clause `schedule` controls how to split iterations and assign them to threads
 - `#pragma omp parallel for schedule(static, chunk_size)`
split iteration space in fixed blocks of `chunk_size`, assign them to threads in round-robin fashion
 - `#pragma omp parallel for schedule(dynamic, chunk_size)`
split iteration space in fixed blocks of `chunk_size`, assign a block to a thread as it becomes idle
 - `#pragma omp parallel for schedule(guided, chunk_size)`
same as `dynamic`, but exponentially reduces the size of the chunk to reduce idling

OpenMP Runtime Library

- A set of routines that deal with:
 - getting info about environment (e.g. `omp_get_num_procs()`)
 - controlling thread creation
 - controlling mutual exclusion
 - timing portions of code
 - `omp_get_wtime()` : elapsed wall-clock time (in seconds) since arbitrary reference time

Compilation

- You must add special flags to your compilation commands
 - `-fopenmp` for the compilation
 - `-lgomp` for the linker

Note for Mac OS X users

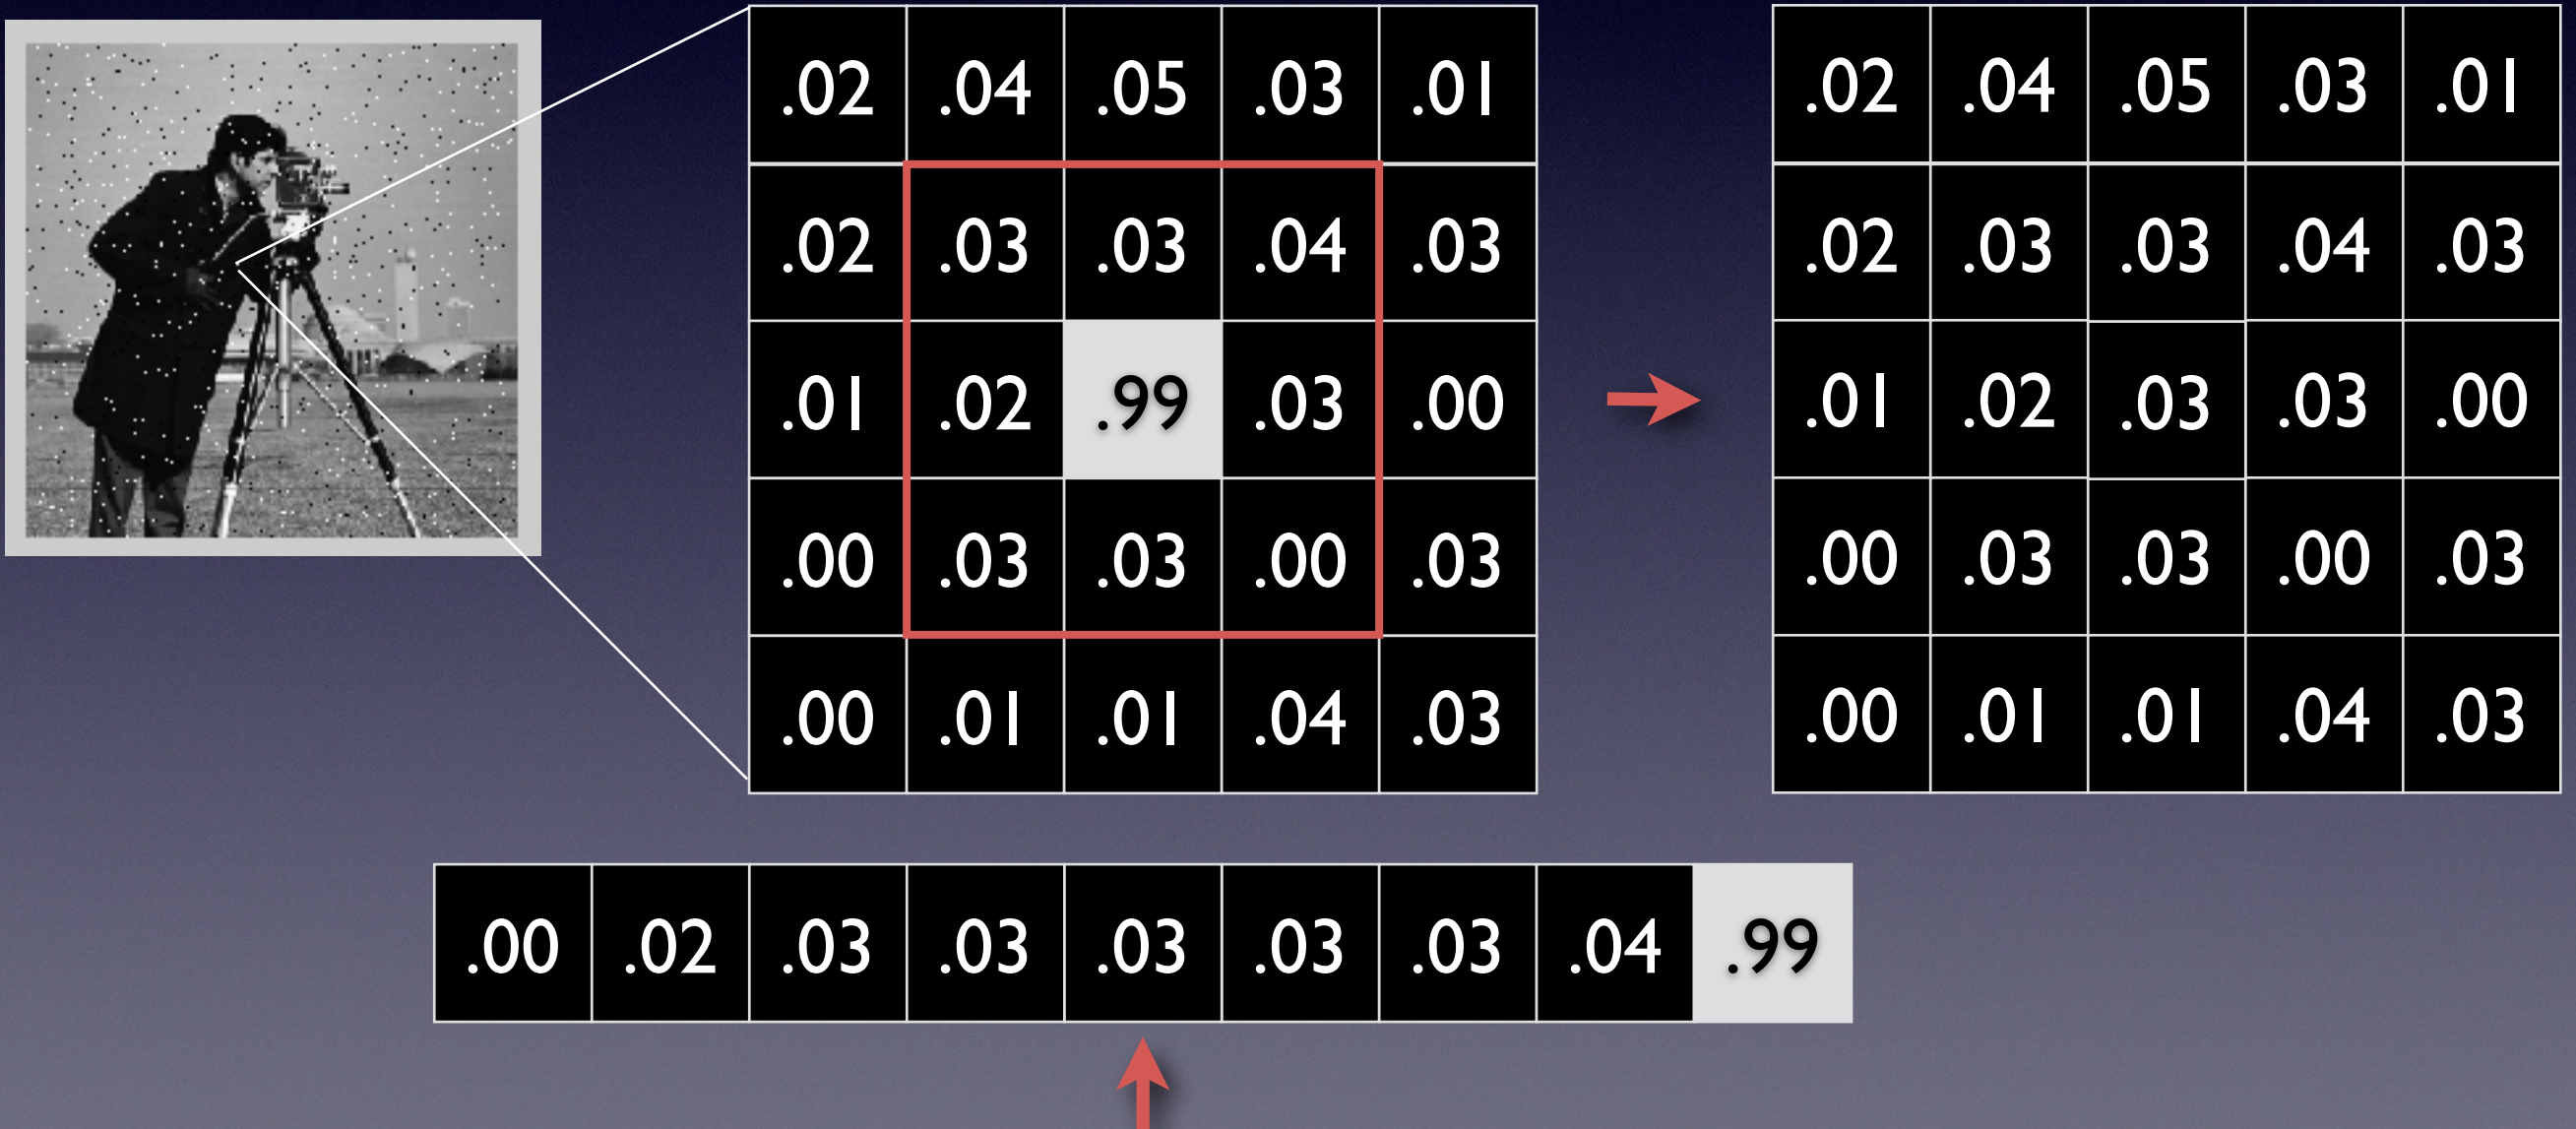
- The version of clang provided does not support OpenMP!
 - <http://openmp.org/wp/openmp-compilers/>
- Solutions:
 - get gcc on Mac OS :-)
 - use homebrew
 - <https://solarianprogrammer.com/2013/06/11/compiling-gcc-mac-os-x/>
 - <http://hpc.sourceforge.net/>
 - use Linux (maybe virtualized, using VirtualBox)
 - ... (but make sure it works!)
- Contact me if you need assistance

Exercise 4

Parallel Median Filtering
of Image Sets using OpenMP


Median Filtering on Image Sets

- We are already familiar with median filtering of images...




Parallel Median Filtering of Images

- In this exercise we'll consider *multiple input image*
 - each to be filtered with a filter of the same size
- Focus of the exercise is on parallelisation of the computation
 - we'll consider two strategies



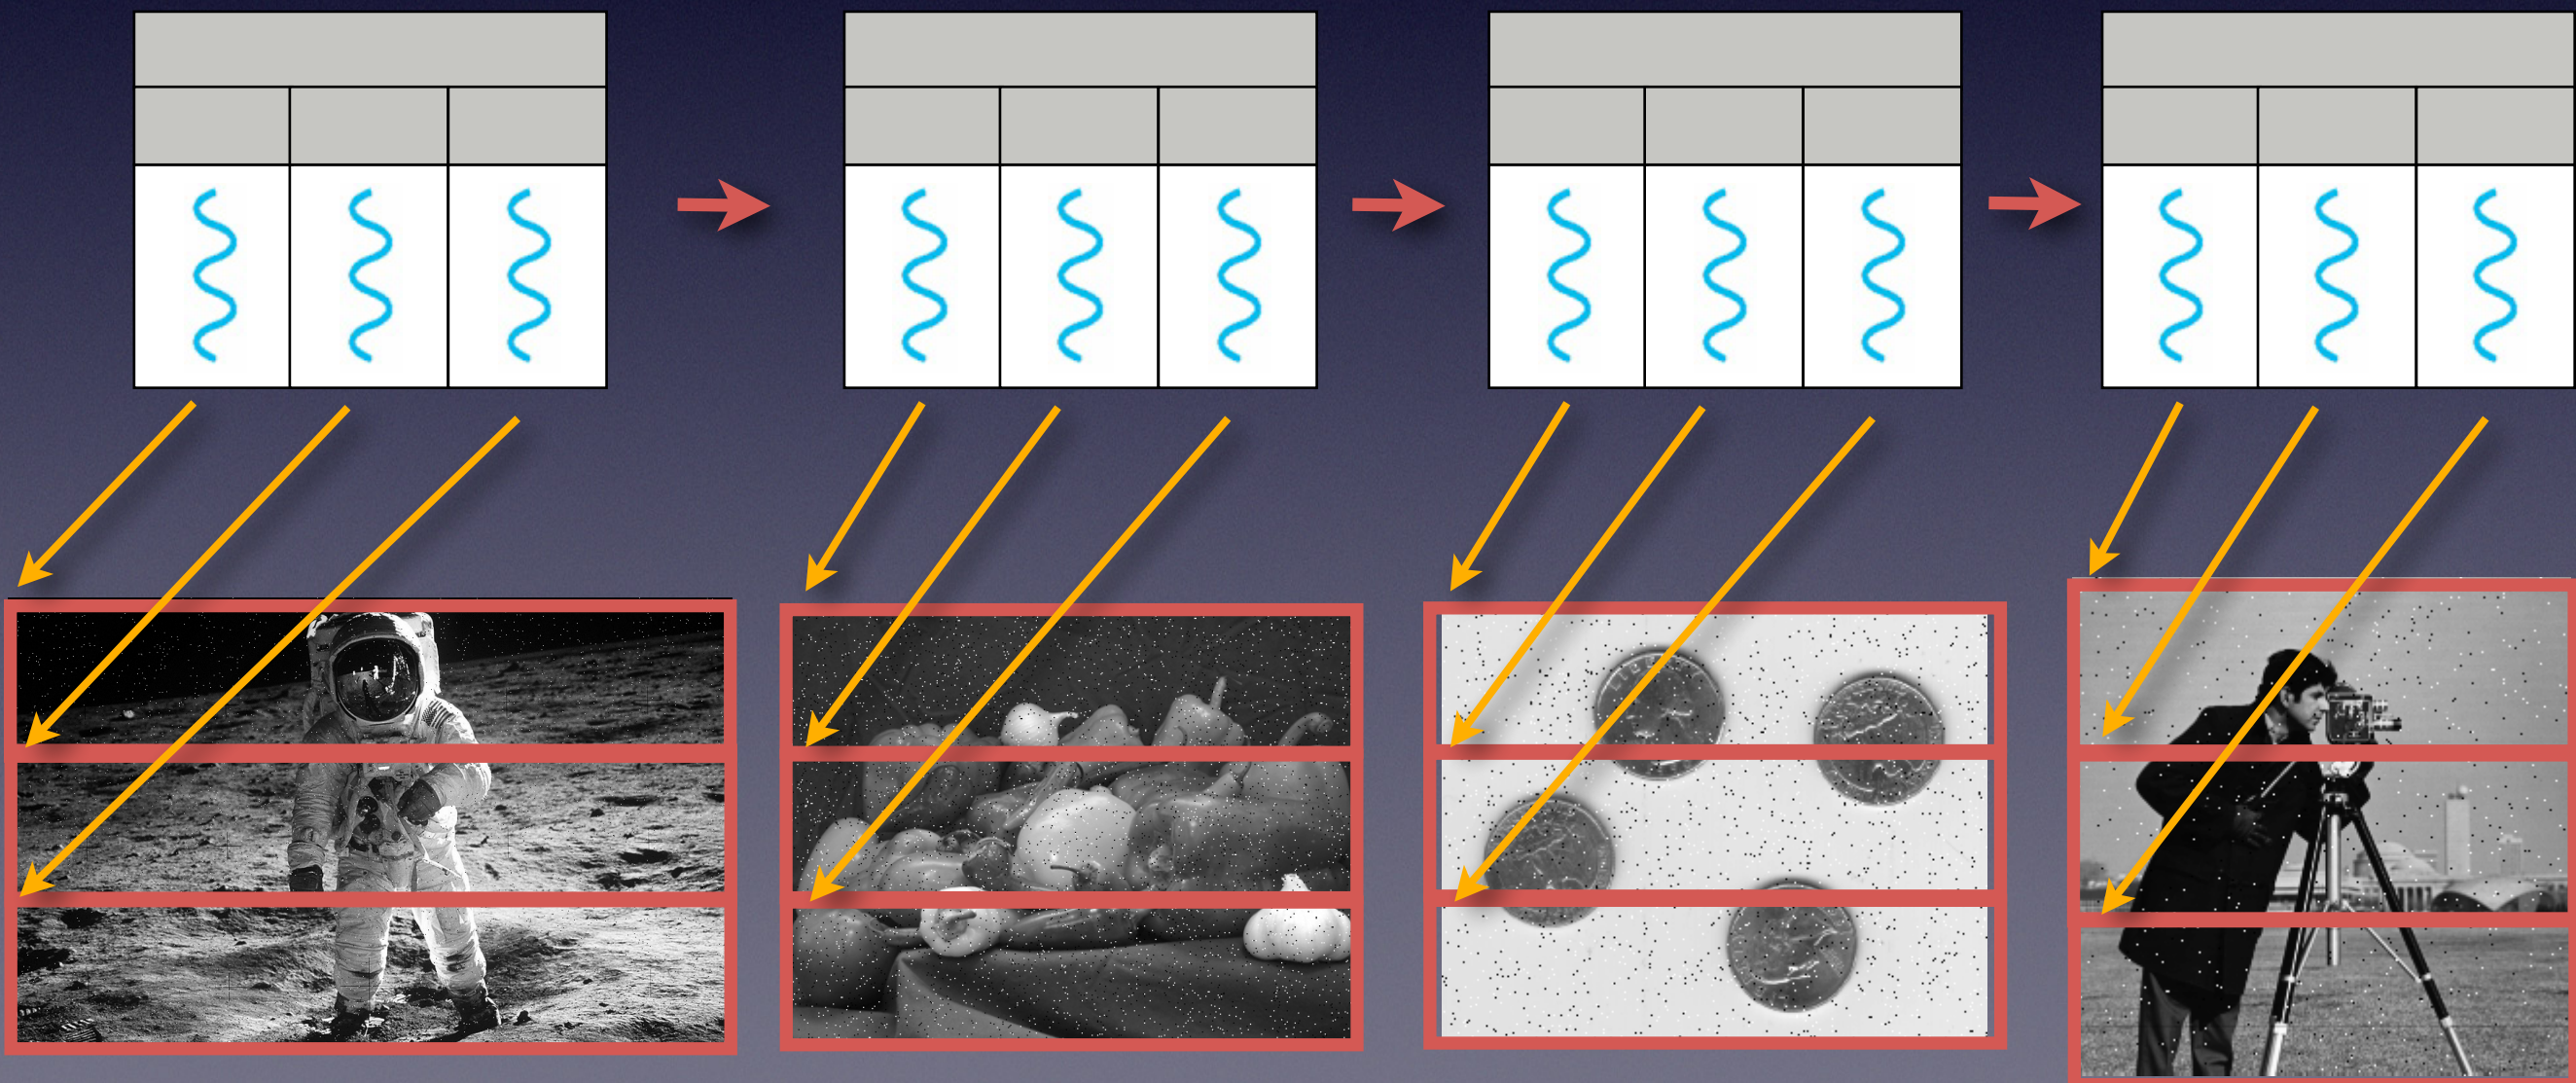
Strategy I: Image-level Parallelism

- Create multiple threads, each processes one or more images
 - each image is processed *serially*, one pixel after another
- Set of input images split among available threads



Strategy 2: Pixel-level Parallelism

- Consider each image one after another
- For each image, filter the pixels in parallel, creating multiple threads and assigning each a sub-set of pixels



Comparing the approaches

- How fast is the parallel version w.r.t. to a serial execution?
Which parallelisation strategy works better?
- Implement a *benchmark* mode
 - on the same set of images, run sequentially the serial version of the computation, the one parallelised at image-level and the one parallelised at pixel level
 - keep track of the computation time for each mode

Technical Notes

- The parallelisation should be achieved using OpenMP
 - mainly `#pragma omp parallel` for directives
 - use `omp_get_wtime()` to track the processing time
- The number of threads to be created is passed as command-line argument
 - 2 main choices with OpenMP: clause after `parallel` for directive or function in runtime library

Additional Details

- Input images stored in text files
 - same format as for Ex. 3
 - filenames are provided as command-line argument
- Window size of filter provided as command-line argument
- Your application should be able to run in one of 4 modes:
 - serial (no parallelism at all)
 - parallel, using strategy 1 (image-level parallelism)
 - parallel, using strategy 2 (pixel-level parallelism)
 - benchmark
 - a command-line argument to select which mode to run
- Number of threads: also command-line argument

Additional Details - 2

- Unless running benchmark mode, write filtered images to files
 - filename for each filtered image: OUT_ + name of text file containing original image
- In benchmark mode print a summary of the computation time spent in each mode when *all* processing is done (i.e., when the 3 modes have completed)