

# Systems Software HS15

## Lab Exercises

Claudio Mura  
[claudio@ifi.uzh.ch](mailto:claudio@ifi.uzh.ch)



# Practical notes on PThreads Mutexes



# Mutual Exclusion with Pthreads

- Main tool: *mutex lock*
  - two states: locked or unlocked
  - locking/unlocking operations are atomic
- Pthread implementation: `pthread_mutex_t`
  - initialization: `pthread_mutex_init( attr )`
    - alternative: `pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;`
  - lock/unlock operations:  
`pthread_mutex_lock()` / `pthread_mutex_unlock()`
  - unlock must be called by the thread currently owning the lock!
- Note: `pthread_mutex_lock()` is blocking!
  - alternative: `pthread_mutex_trylock()`
  - returns immediately if mutex is locked




# Exercise 6

Simulating the Behaviour of a  
Memory Allocator in a  
Multi-process Environment



# Our Toy Environment

- Available memory divided into  $B$  blocks of fixed size
  - managed by a memory allocator
  - a block can be unallocated or assigned to a process
- $P$  processes running on the system
  - need some extra memory blocks to perform computation
  - request blocks to the memory allocator, release them when done






# Behaviour of a Process

- Every  $T$  msec performs some compute-intensive task
  - requests blocks (no. depends on the task) from the allocator
  - if allocation ok, spend some time computing, then release blocks
  - if allocation fails (no blocks available), abort and do nothing
- Spends rest of the time sleeping



# Behaviour of the Allocator

- Keeps a list with information on each block
  - allocation status: free (-1) or allocated to a process (id of process)
- Answers allocation requests of processes
  - check if the number of blocks requested is available
  - if yes, assigns blocks and marks them; if not, does nothing; in both cases, notifies the process of the outcome



Proc. 1

Proc. 2

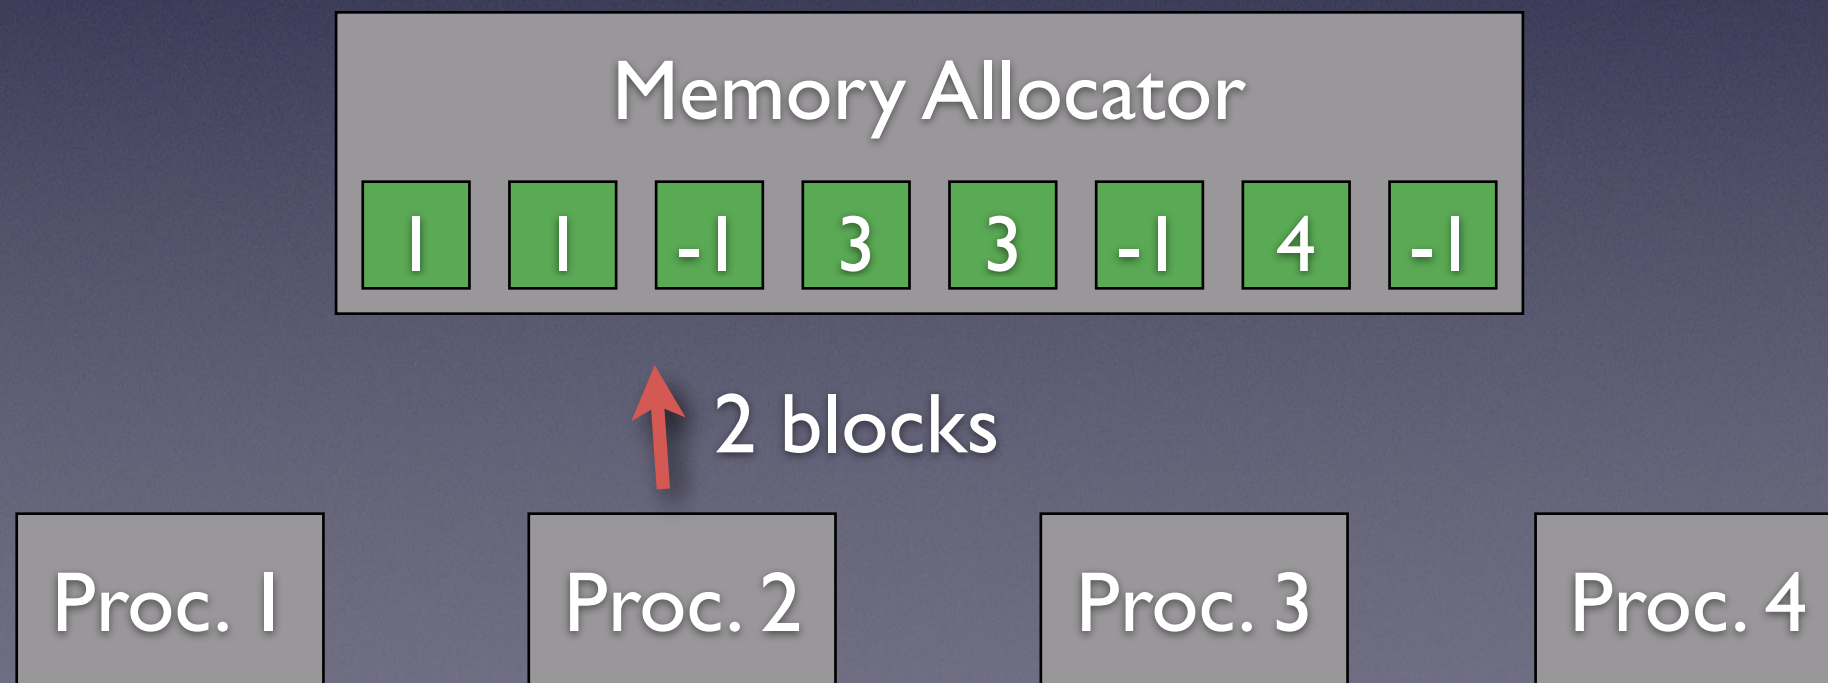
Proc. 3

Proc. 4



# Behaviour of the Allocator


- Keeps a list with information on each block
  - allocation status: free (-1) or allocated to a process (id of process)
- Answers allocation requests of processes
  - check if the number of blocks requested is available
  - if yes, assigns blocks and marks them; if not, does nothing; in both cases, notifies the process of the outcome





# Behaviour of the Allocator


- Keeps a list with information on each block
  - allocation status: free (-1) or allocated to a process (id of process)
- Answers allocation requests of processes
  - check if the number of blocks requested is available
  - if yes, assigns blocks and marks them; if not, does nothing; in both cases, notifies the process of the outcome





# Behaviour of the Allocator

- Keeps a list with information on each block
  - allocation status: free (-1) or allocated to a process (id of process)
- Answers allocation requests of processes
  - check if the number of blocks requested is available
  - if yes, assigns blocks and marks them; if not, does nothing; in both cases, notifies the process of the outcome



Proc. 1

Proc. 2

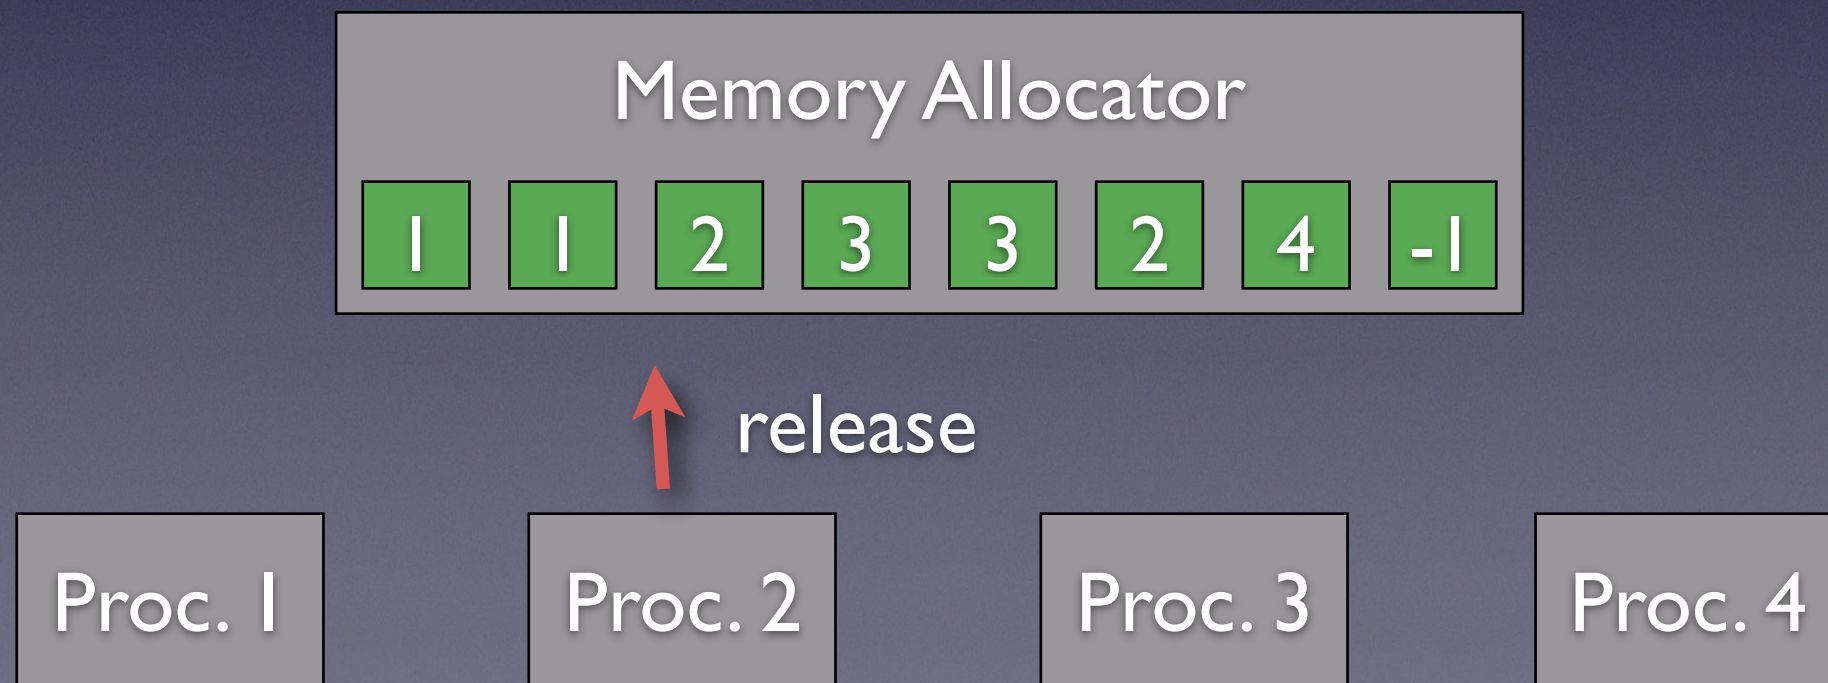
Proc. 3

Proc. 4



# Behaviour of the Allocator

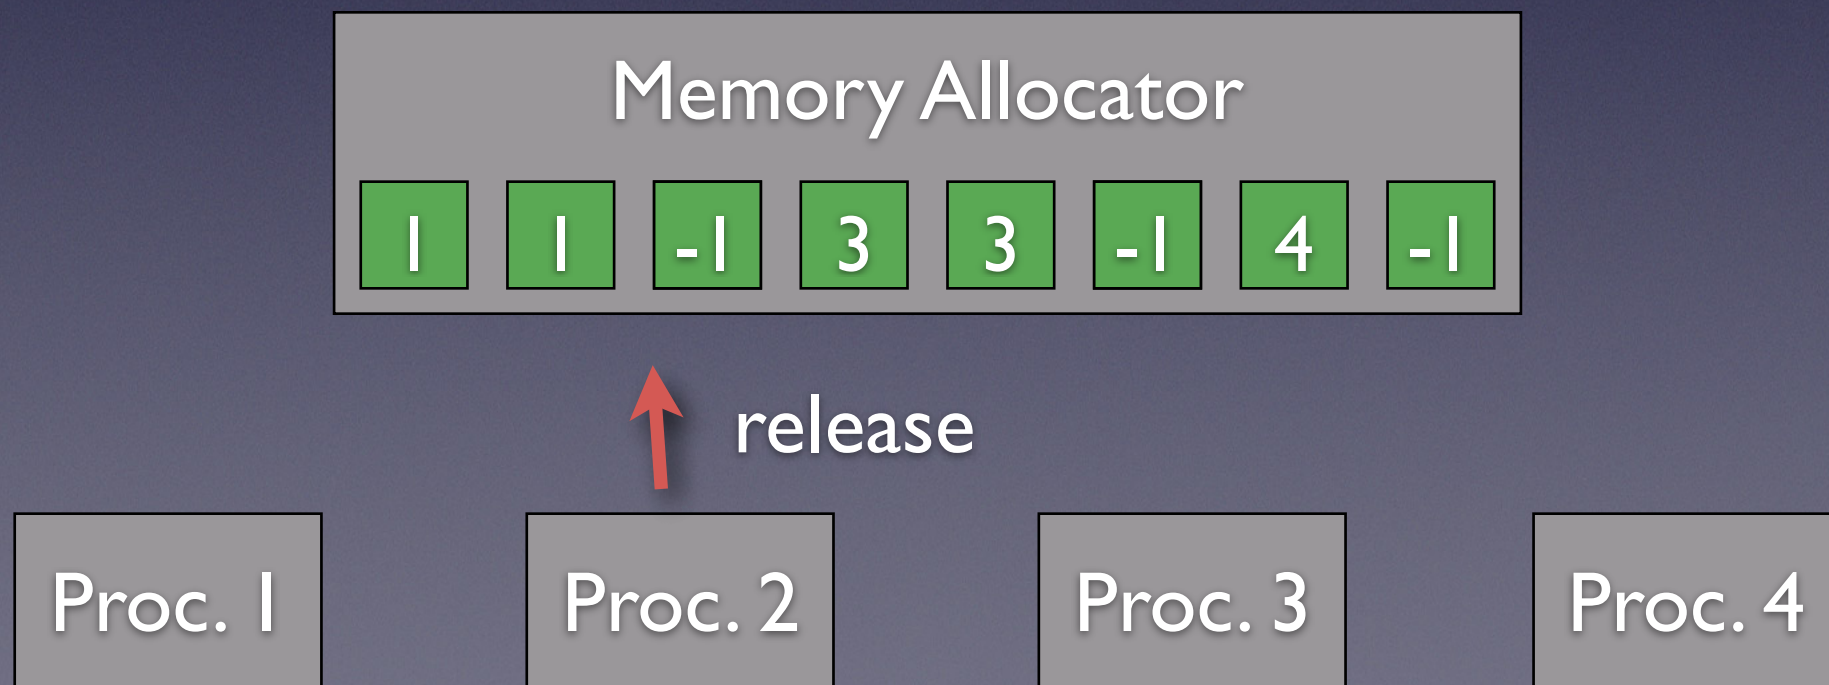
- Keeps a list with information on each block
  - allocation status: free (-1) or allocated to a process (id of process)
- Answers allocation requests of processes
  - check if the number of blocks requested is available
  - if yes, assigns blocks and marks them; if not, does nothing; in both cases, notifies the process of the outcome





# Behaviour of the Allocator

- Keeps a list with information on each block
  - allocation status: free (-1) or allocated to a process (id of process)
- Answers allocation requests of processes
  - check if the number of blocks requested is available
  - if yes, assigns blocks and marks them; if not, does nothing; in both cases, notifies the process of the outcome





# Simulation

- Implementing behaviour of processes:
  - use one thread (PThreads) for every process
    - each thread identified by an id
  - sleep time after work (attempt):  $T$  ms,  $T$  random in  $[T_{min}, T_{max}]$
  - number of blocks needed:  $W$ , random in  $[W_{min}, W_{max}]$
  - simulate computing: sleep for  $T'$  ms,  $T'$  random in  $[T_{min}, W \times T_{max}]$

**do**

generate random number  $W$  in the range  $[W_{min}, W_{max}]$

request allocation of  $W$  blocks of memory

**if** allocation successful

sleep for some random time  $T'$  in the range  $[T_{min}, W \times T_{max}]$

release allocated blocks

**end**

sleep for some random time  $T$  in the range  $[T_{min}, T_{max}]$

**while** simulation is running



# Simulation

- Implementing behaviour of allocator:
  - implement the list of block entries as `std::vector< int >`
    - each entry: -1 (block not allocated) or id of thread using the block
  - provide 2 functions/methods to allocate and release memory blocks
  - ensure that accesses to the vector are properly synchronized!
    - use a `pthread_mutex_t` (and nothing else!)



# Additional Details

- Command line arguments (in order):
  - # processes  $P$  in the simulated system (i.e. threads to be created)
  - # memory blocks  $B$  (i.e. size of the `std::vector< int >`)
  - values  $W_{min}$ ,  $W_{max}$ ,  $T_{min}$ ,  $T_{max}$
- Making threads sleep: use `nanosleep()` function
- Provide console output to show how simulation evolves
  - warning: terminal is a shared resource! use mutex
- Terminate simulation when user presses 'e' + return
  - make sure threads terminate gracefully